

# How to determine the optimal number of clusters for k-means clustering using Elbow Method

## Introduction

K-means is a type of unsupervised learning and one of the popular methods of clustering unlabelled data into k clusters. One of the trickier tasks in clustering is identifying the appropriate number of clusters k. In this tutorial, we will provide an overview of how k-means works and discuss how to implement your own clusters.

We will also understand how to use the elbow method as a way to estimate the value k. Another popular method of estimating k is through silhouette analysis, a scikit learn example can be found [here](#).

We will use the wholesale customer dataset which can be downloaded [here](#).

## K-means Overview

Before diving into the dataset, let us briefly discuss how k-means works:

1. The process begins with k centroids initialised at random.
2. These centroids are used to assign points to its nearest cluster.
3. The mean of all points within the cluster is then used to update the position of the centroids.
4. The above steps are repeated until the values of the centroids stabilise.

## Getting Started

In this tutorial, we will be using the scikit-learn's implementation of k-means which can be found [here](#)

## The dataset

The dataset we will study refers to clients of a wholesale distributor. It contains information such as clients annual spend on fresh product, milk products, grocery products etc. Below is some more information an each feature:

1. FRESH: annual spending (m.u.) on fresh products (Continuous)

2. MILK: annual spending (m.u.) on milk products (Continuous)
3. GROCERY: annual spending (m.u.) on grocery products (Continuous)
4. FROZEN: annual spending (m.u.) on frozen products (Continuous)
5. DETERGENTS\_PAPER: annual spending (m.u.) on detergents and paper products (Continuous)
6. DELICATESSEN: annual spending (m.u.) on delicatessen products (Continuous)
7. CHANNEL: customer channels - Horeca (Hotel/Restaurant/Cafe) or Retail channel (Nominal)
8. REGION: customer regions - Lisbon, Oporto or Other (Nominal)

```
# Import required packages
import pandas as pd
from sklearn.preprocessing import MinMaxScaler
from sklearn.cluster import KMeans

import matplotlib.pyplot as plt
```

Read in data and inspect the first 5 records.

```
data = pd.read_csv('Wholesale customers data.csv')
data.head()
```

|   | Channel | Region | Fresh | Milk | Grocery | Frozen | Detergents_Paper | Delicassen |
|---|---------|--------|-------|------|---------|--------|------------------|------------|
| 0 | 2       | 3      | 12669 | 9656 | 7561    | 214    | 2674             | 1338       |
| 1 | 2       | 3      | 7057  | 9810 | 9568    | 1762   | 3293             | 1776       |
| 2 | 2       | 3      | 6353  | 8808 | 7684    | 2405   | 3516             | 7844       |
| 3 | 1       | 3      | 13265 | 1196 | 4221    | 6404   | 507              | 1788       |
| 4 | 2       | 3      | 22615 | 5410 | 7198    | 3915   | 1777             | 5185       |

Below is a split of categorical and continuous features

```
categorical_features = ['Channel', 'Region']
continuous_features = ['Fresh', 'Milk', 'Grocery', 'Frozen',
'Detergents_Paper', 'Delicassen']
```

Descriptive statistics below shows on average clients spend the most on fresh groceries and the least on delicassen.

```
data[continuous_features].describe()
```

|              | Fresh         | Milk         | Grocery      | Frozen       | Detergents_Paper | Delicassen   |
|--------------|---------------|--------------|--------------|--------------|------------------|--------------|
| <b>count</b> | 440.000000    | 440.000000   | 440.000000   | 440.000000   | 440.000000       | 440.000000   |
| <b>mean</b>  | 12000.297727  | 5796.265909  | 7951.277273  | 3071.931818  | 2881.493182      | 1524.870455  |
| <b>std</b>   | 12647.328865  | 7380.377175  | 9503.162829  | 4854.673333  | 4767.854448      | 2820.105937  |
| <b>min</b>   | 3.000000      | 55.000000    | 3.000000     | 25.000000    | 3.000000         | 3.000000     |
| <b>25%</b>   | 3127.750000   | 1533.000000  | 2153.000000  | 742.250000   | 256.750000       | 408.250000   |
| <b>50%</b>   | 8504.000000   | 3627.000000  | 4755.500000  | 1526.000000  | 816.500000       | 965.500000   |
| <b>75%</b>   | 16933.750000  | 7190.250000  | 10655.750000 | 3554.250000  | 3922.000000      | 1820.250000  |
| <b>max</b>   | 112151.000000 | 73498.000000 | 92780.000000 | 60869.000000 | 40827.000000     | 47943.000000 |

To use the categorical features, we need to convert the categorical features to binary using pandas get dummies.

```
for col in categorical_features:
    dummies = pd.get_dummies(data[col], prefix=col)
    data = pd.concat([data, dummies], axis=1)
    data.drop(col, axis=1, inplace=True)

data.head()
```

|          | Fresh | Milk | Grocery | Frozen | Detergents_Paper | Delicassen | Channel_1 | Channel_2 | Region_1 | Region_2 | Region_3 |
|----------|-------|------|---------|--------|------------------|------------|-----------|-----------|----------|----------|----------|
| <b>0</b> | 12669 | 9656 | 7561    | 214    | 2674             | 1338       | 0         | 1         | 0        | 0        | 1        |
| <b>1</b> | 7057  | 9810 | 9568    | 1762   | 3293             | 1776       | 0         | 1         | 0        | 0        | 1        |
| <b>2</b> | 6353  | 8808 | 7684    | 2405   | 3516             | 7844       | 0         | 1         | 0        | 0        | 1        |
| <b>3</b> | 13265 | 1196 | 4221    | 6404   | 507              | 1788       | 1         | 0         | 0        | 0        | 1        |
| <b>4</b> | 22615 | 5410 | 7198    | 3915   | 1777             | 5185       | 0         | 1         | 0        | 0        | 1        |

To give equal importance to all features, we need to scale the continuous features. We will be using scikit-learn's MinMaxScaler as the feature matrix is a mix of binary and continuous features . Other alternatives includes StandardScaler.

```
mms = MinMaxScaler()
mms.fit(data)
data_transformed = mms.transform(data)
```

For each k value, we will initialise k-means and use the inertia attribute to identify the sum of squared distances of samples to the nearest cluster centre.


```
Sum_of_squared_distances = []
K = range(1,15)
for k in K:
```

```
km = KMeans(n_clusters=k)
km = km.fit(data_transformed)
Sum_of_squared_distances.append(km.inertia_)
```

As  $k$  increases, the sum of squared distance tends to zero. Imagine we set  $k$  to its maximum value  $n$  (where  $n$  is number of samples) each sample will form its own cluster meaning sum of squared distances equals zero.

Below is a plot of sum of squared distances for  $k$  in the range specified above. If the plot looks like an arm, then the elbow on the arm is optimal  $k$ .

```
plt.plot(K, Sum_of_squared_distances, 'bx-')
plt.xlabel('k')
plt.ylabel('Sum_of_squared_distances')
plt.title('Elbow Method For Optimal k')
plt.show()
```



In the plot above the elbow is at  $k=5$  indicating the optimal  $k$  for this dataset is around 5. Therefore, you should review the results for 4, 5 and 6 while designing your  $k$ -means algorithm.

Source